

Citation Integrity Audit

Paper audited: Context-Window Degradation and Its Impact on Long-Document Research Synthesis

Audit method: Every reference in the paper's "Complete References" list was checked against arXiv, the publisher/DOI record, or (for one software resource) the GitHub repository itself, for: correct title, correct authors, correct year, correct venue, a working DOI/arXiv ID, and confirmation that the paper genuinely exists and matches the citing text. Verification sources used: arXiv.org, Semantic Scholar, publisher DOI pages, ACL Anthology, and OpenReview/ICLR proceedings pages.

Verification Log — In-Paper References

#	Reference (short form)	Title/Authors/ Year	Venue	Link Resolves	Notes
1	Bai et al., 2024 — LongBench	Verified	ACL 2024	arXiv:2308.14508	Confirmed via GitHub (THUDM/LongBench) and ACL Anthology.
2	Hsieh et al., 2024 — RULER	Verified	COLM 2024	arXiv:2404.06654	Confirmed via official repo hsiehjackson/RULER.
3	Huang et al., 2023 — Survey	Verified	arXiv preprint	arXiv:2311.12351	Correctly presented as an unreviewed preprint.
4	Kamradt, 2023 — NIAH	Verified	N/A (software)	GitHub repo	Correctly cited as software, not a paper.
5	Kazemnejad et al., 2023	Verified	NeurIPS 2023	arXiv:2305.19466	Matches Section 2.1 claim on PE generalization.
6	Li, K. et al., 2025 — LaRA	Verified	arXiv preprint	arXiv:2502.09977	Confirmed critique of prior RAG-vs-LC benchmark design.
7	Li, Z. et al., 2024 — RAG or LC?	Verified	EMNLP 2024 Industry	arXiv:2407.16833	Self-Route method correctly attributed.
8	Liu, J. et al., 2025 — Survey	Verified	arXiv preprint	arXiv:2503.17407	Large author list abbreviated appropriately.
9	Liu, N. F. et al., 2024 — Lost in the Middle	Verified	TACL, Vol. 12	DOI 10.1162/tacl_a_00638	Core citation; U-shaped curve finding matches exactly.
10	Press et al., 2022 — ALiBi	Verified	ICLR 2022	arXiv:2108.12409	Confirmed.
11	Su et al., 2024 — RoFormer	Verified	Neurocomputing	DOI 10.1016/j.neucom.2023.127063	Journal (2024) version correctly cited over 2021 preprint.
12	Xiao et al., 2024 — StreamingLLM	Verified	ICLR 2024	arXiv:2309.17453	Confirmed via official repo mit-han-lab/streaming-llm.
13	Xu et al., 2024 — Retrieval meets LC	Verified	ICLR 2024	arXiv:2310.03025	In reference list; not explicitly named in body text (flagged below).

Result: 13/13 references verified as genuine, correctly attributed, and correctly linked. No fabricated, mismatched, or non-existent references were found.

Observations / Minor Issues Flagged

- Xu et al. (2024)** appears in the reference list but is not explicitly named in the body text (it may be folded into general discussion of RAG/long-context hybrids in Section 2.4/4.1). Recommendation: either cite it explicitly where relevant or remove it — currently it is an orphaned reference. It was kept in the companion repository's collection because it is independently verifiable and topically relevant.
- In-text claims were spot-checked against source abstracts, not full-text re-derivation of every statistic (e.g., specific accuracy percentages from RULER or Liu et al. were not independently re-run). This audit confirms the references exist, are correctly attributed, and plausibly support the claims made — not that every quantitative figure was independently reproduced, which would require re-running the cited benchmarks and is outside the scope of a citation-integrity audit.
- No AI-generated or “hallucinated” citations were detected. All DOIs and arXiv IDs resolved to the correct paper on the first attempt.

Additional Papers Added for the Companion Repository

To meet the 20-paper minimum for the Awesome repository (the review paper itself cites 13 sources), 7 additional papers were sourced and independently verified using the same method before being added to the repository's references list:

- Hsieh, C.-Y. et al. (2024) — *Found in the Middle* (arXiv:2406.16008)
- Chen, S. et al. (2023) — *Position Interpolation* (arXiv:2306.15595)
- Peng, B. et al. (2024) — *YaRN* (arXiv:2309.00071)
- Jiang, H. et al. (2023) — *LLMLingua* (arXiv:2310.05736)
- Jiang, H. et al. (2023) — *LongLLMLingua* (arXiv:2310.06839)
- Pan, Z. et al. (2024) — *LLMLingua-2* (arXiv:2403.12968)
- Jiang, Z. et al. (2024) — *LongRAG* (arXiv:2406.15319)
- Ji, Z. et al. (2023) — *Survey of Hallucination in NLG* (arXiv:2202.03629)

Each was checked for title/author/year accuracy and link resolution against arXiv directly, following the same verification standard applied above.