Context-Window Degradation and Its Impact on Long-Document Research Synthesis

A Review of Positional, Attentional, and Architectural Constraints on Large Language Model Performance over Extended Inputs

Abstract

Large language models (LLMs) now advertise context windows spanning hundreds of thousands to millions of tokens, prompting expectations that entire corpora, case files, or literature sets can be synthesized in a single pass. Empirical evidence, however, indicates that nominal context length and effective context utilization diverge substantially. This paper reviews the phenomenon of context-window degradation—the systematic decline in a model's ability to locate, weight, and integrate information as input length grows or as relevant content is displaced from the beginning or end of the input—and examines its consequences for long-document research synthesis, a task that requires uniform attention across dozens or hundreds of source passages regardless of their position. Drawing on positional-bias studies, synthetic and realistic long-context benchmarks, and comparative evaluations of retrieval-augmented and long-context architectures, the paper characterizes the primary mechanisms of degradation: primacy-recency ("lost-in-the-middle") effects, positional-encoding extrapolation failure, attention dilution, and distractor accumulation. It then surveys current mitigation strategies, including retrieval-augmented generation, hybrid routing, prompt compression, attention-sink and sparse-attention mechanisms, and calibration techniques, while assessing their trade-offs in cost, faithfulness, and coverage. The review identifies persistent gaps—the absence of standardized synthesis-specific benchmarks, limited understanding of multi-hop aggregation failures, and inadequate faithfulness metrics for long-form synthesis—and proposes directions for future research, including position-invariant evaluation protocols and hybrid architectures tailored to research-synthesis workloads.

Keywords: *context window; long-context language models; lost-in-the-middle; positional bias; retrieval-augmented generation; attention degradation; research synthesis; large language models; benchmark evaluation*

1. Introduction

Large language models have undergone a rapid expansion in nominal context capacity, from the 2,048-token windows typical of early transformer-based systems to advertised windows of 128,000, 200,000, or over one million tokens in current commercial and open-weight models. This expansion has been marketed as enabling a new class of applications: feeding an LLM an entire book, a full set of legal filings, or a corpus of dozens of scientific papers, and asking it to synthesize a coherent answer without external retrieval or chunking. For research synthesis in particular—the task of aggregating, comparing, and reconciling claims distributed across many long source documents—this promise is attractive, since it would appear to eliminate the information loss inherent in retrieval-based chunking pipelines.

A growing body of empirical work complicates this picture. Liu et al. (2024) demonstrated that language models exhibit a pronounced U-shaped performance curve when the information required to answer a question is repositioned within a long input: accuracy is highest when the relevant passage sits at the very beginning or end of the context and degrades substantially when it is placed in the middle, a pattern now widely known as the "lost-in-the-middle" effect. Subsequent benchmarks have shown that this positional sensitivity is only one facet of a broader phenomenon in which nominal and effective context length diverge as input grows (Hsieh et al., 2024). For tasks such as long-document research synthesis, which require a model to attend not to a single needle but to many dispersed, position-independent facts and to integrate them faithfully, these degradation effects pose a direct threat to output reliability.

This paper reviews the mechanisms, evidence, and mitigations associated with context-window degradation, with specific attention to its implications for long-document research synthesis. Section 2 situates the problem within the broader literature on transformer architecture, positional encoding, and long-context evaluation. Section 3 provides a technical account of the principal degradation mechanisms. Section 4 surveys current mitigation approaches, including retrieval-augmented generation (RAG), hybrid routing, and attention-management techniques. Section 5 discusses methodological and practical limitations of the existing evidence base. Section 6 identifies research gaps and proposes directions for future work, and Section 7 concludes.

2. Background and Related Work

2.1 The Transformer Attention Mechanism and Its Scaling Limits

The self-attention mechanism at the core of transformer architectures computes pairwise interactions between all tokens in an input sequence, giving these models in principle the capacity to relate any token to any other regardless of distance. In practice, this capacity is constrained by two factors: the computational and memory cost of attention, which scales quadratically with sequence length, and the way positional information is encoded and generalizes beyond the lengths seen during training. Positional encoding schemes—including the original sinusoidal absolute encodings, learned absolute embeddings, relative position encodings, Rotary Position Embedding (RoPE) (Su et al., 2024), and Attention with Linear Biases (ALiBi) (Press et al., 2022)—were each developed in part to improve a model's ability to generalize, or extrapolate, to sequence lengths beyond its training distribution. Kazemnejad et al. (2023) conducted a systematic comparison of these schemes and found that even commonly used relative encodings such as RoPE and ALiBi provide only partial protection against length-generalization failure, and that no explicit positional encoding scheme reliably matches the length-generalization performance achieved implicitly by models trained without any explicit positional signal on certain reasoning tasks. This body of work establishes that context-window expansion is not merely an engineering matter of allocating more memory, but is bound up with unresolved questions about how position and content are jointly represented.

2.2 Evidence of Positional Bias: Lost in the Middle

The most widely cited empirical demonstration of context degradation is the study by Liu et al. (2024), who evaluated language models on multi-document question answering and synthetic key-value retrieval tasks while systematically varying the position of the relevant document or key-value pair within the input. They found that performance was highest when the relevant content occurred at the beginning or end of the context and fell substantially, in some cases below the performance obtained with no relevant document at all, when the content was placed in the middle. Critically, they showed that this effect persisted even in models explicitly designed and marketed for long-context use, and that extending a model's context window did not by itself resolve the bias. This finding has been replicated and extended across model families and tasks; a survey of robustness limitations characterizes the resulting performance pattern as a "characteristic U-shaped curve" that recurs across both retrieval and reasoning-oriented long-context benchmarks. Subsequent work examining multi-hop question answering found that the severity of degradation depends not only on the absolute position of relevant information but also on the relative distance between multiple pieces of information that must be combined, indicating that the effect compounds when synthesis requires integrating several dispersed facts rather than retrieving a single fact.

2.3 Beyond Needle-in-a-Haystack: Comprehensive Long-Context Benchmarks

Early diagnostics of long-context capability relied heavily on the "needle-in-a-haystack" (NIAH) paradigm introduced by Kamradt (2023), in which a single fact is inserted at a controlled depth within a long distractor text and the model is asked to retrieve it. While influential, this test has been criticized as measuring only a superficial form of long-context competence, since near-perfect NIAH scores do not guarantee robust performance on tasks requiring reasoning, aggregation, or tracking of multiple variables. Hsieh et al. (2024) addressed this limitation with RULER, a synthetic benchmark that extends the NIAH paradigm to include multiple needles, multi-hop tracing, and aggregation tasks, and found that despite near-perfect vanilla NIAH accuracy, most evaluated long-context models exhibited substantial performance drops as input length increased, such that a model's effective context length was often far shorter than its claimed maximum. LongBench (Bai et al., 2024) complements this synthetic approach with a bilingual, multi-task benchmark built from realistic long documents spanning single- and multi-document question answering, summarization, and few-shot learning, and similarly found that even strong commercial models struggled as document length increased. Together, these benchmarks establish that context degradation is a general property of current architectures rather than an artifact of a particular synthetic task.

2.4 Retrieval-Augmented Generation as an Alternative Paradigm

Retrieval-augmented generation (RAG) offers an architecturally distinct response to the long-context problem: rather than presenting the entire corpus to the model, a retriever first selects a small set of relevant passages, which are then supplied as a much shorter context. Comparative studies of RAG and long-context (LC) processing have produced genuinely mixed findings depending on task and evaluation design. Li et al. (2024) found that, when sufficiently resourced, long-context processing consistently outperformed RAG in average accuracy across several public datasets, though RAG retained a substantial cost advantage, motivating their proposed hybrid "Self-Route" method that dynamically routes queries to RAG or long-context processing based on model self-reflection. Other work reaches different conclusions: Li et al. (2025), introducing the LaRA benchmark, argue that prior comparisons suffer from benchmark-design flaws and find that RAG performs comparably to long-context processing on single-location retrieval tasks and offers a distinct advantage in identifying unanswerable or hallucination-prone queries, whereas long-context processing retains an edge on reasoning and cross-document comparison tasks.

This inconsistency across studies is itself evidence that context degradation interacts heavily with task type, and that no single architecture is uniformly superior for synthesis workloads.

3. Main Technical Discussion: Mechanisms of Context-Window Degradation

Context-window degradation is not a single phenomenon but the aggregate effect of several interacting mechanisms. This section characterizes four that are most consequential for long-document research synthesis.

3.1 Primacy-Recency Bias and the Lost-in-the-Middle Effect

As established in Section 2.2, transformer-based LLMs disproportionately weight information near the start and end of an input context. For research synthesis, where relevant claims are typically distributed across many source passages with no natural correlation between position and importance, this bias means that the ordering in which documents are concatenated can materially change which facts are incorporated into a synthesized answer. A synthesis system that processes ten source papers, for example, risks systematically underweighting evidence from papers five through eight relative to papers one and ten, independent of the actual strength or relevance of their findings—a distortion with no analogue in careful human literature review, where a reader's attention is not structurally biased by a document's position in a stack.

3.2 Positional-Encoding Extrapolation Failure

A second mechanism concerns what happens when input length exceeds the range seen during training. Because most positional encoding schemes are calibrated on a fixed training-length distribution, extending inference to longer sequences—whether through native long-context training or post hoc extension techniques such as position interpolation—can push position representations into a regime the model has not learned to interpret reliably. Kazemnejad et al. (2023) show that this extrapolation failure is scheme-dependent and task-dependent, and surveys of transformer architecture for long-context modeling catalogue a range of scaling and interpolation strategies developed specifically to counteract this degradation, indicating that it is treated as a distinct engineering problem from attentional bias per se. For synthesis tasks operating near or beyond a model's trained context length, this failure mode compounds positional bias: not only is middle content underweighted, but the model's basic sense of token ordering and relative distance can become unreliable.

3.3 Attention Dilution and the Counting/Aggregation Problem

A third mechanism, distinct from positional bias, is attention dilution: as the number of tokens competing for a fixed attention budget grows, the softmax normalization underlying attention forces the model to distribute a bounded quantity of attention mass across an increasing number of positions, reducing the mass any single relevant token can receive. This effect is closely related to the phenomenon of "attention sinks," in which models learn to direct disproportionate attention toward initial tokens largely because early tokens are visible to all subsequent positions during autoregressive training, rather than because those tokens are semantically important (Xiao et al., 2024). RULER's aggregation and multi-hop tracing tasks, which require a model to combine or count information spread across the input rather than retrieve a single fact, show substantially larger performance drops than simple retrieval tasks as context length grows (Hsieh et al., 2024), consistent with attention dilution compounding rather than merely coexisting with positional bias. For research synthesis, which frequently requires exactly this kind of aggregation—for instance, tallying how many source studies support a given claim, or comparing effect sizes across a set of papers—dilution effects are especially damaging because the task cannot be reduced to single-fact retrieval.

3.4 Distractor Accumulation and Noise Sensitivity

Finally, as context length grows, the ratio of task-relevant to task-irrelevant tokens typically declines, particularly in realistic synthesis settings where source documents contain substantial material that is not directly relevant to a given query. Liu et al. (2024) found that increasing the number of distractor documents in multi-document question answering degraded performance even when the position of the relevant document was held constant, indicating that distractor volume is an independent driver of degradation separate from positional placement. This has direct implications for long-document research synthesis: unlike a curated, retrieval-filtered context, a full-document or full-corpus context necessarily includes substantial non-relevant material—background sections, methodological detail irrelevant to the synthesis question, references—that dilutes the effective signal available to the model even before positional and attentional effects are considered.

4. Current Approaches and Developments

A range of architectural and pipeline-level techniques have been developed to counteract context-window degradation. This section surveys the principal families of approaches relevant to long-document research synthesis.

4.1 Retrieval-Augmented Generation and Hybrid Routing

The most mature mitigation strategy remains retrieval-augmented generation, which sidesteps degradation by limiting the model's effective input to a small, retrieval-filtered set of passages rather than an entire corpus. As discussed in Section 2.4, RAG's advantages are most pronounced on cost and on tasks where relevance is concentrated in a few passages, while its drawbacks include sensitivity to retriever quality and reduced ability to perform cross-document reasoning that depends on passages the retriever did not select. Hybrid approaches attempt to combine the strengths of both paradigms: Li et al. (2024) propose Self-Route, which uses a model's own confidence signals to decide, on a per-query basis, whether to answer from retrieved passages or from the full long context, achieving performance comparable to long-context processing at substantially lower token cost. LongRAG similarly restructures the retrieval unit itself, using long retrieval units processed by long-context readers rather than the short passages typical of classical RAG pipelines, aiming to reduce the loss of contextual information introduced by aggressive chunking. For research synthesis specifically, hybrid and adaptive routing strategies are attractive because synthesis queries vary widely in how many source documents are truly relevant, making a fixed retrieval-versus-full-context policy suboptimal across a typical workload.

4.2 Architectural Interventions: Attention Sinks and Sparse Attention

A second family of approaches modifies the attention mechanism itself rather than the surrounding pipeline. StreamingLLM (Xiao et al., 2024) identified the attention-sink phenomenon described in Section 3.3 and showed that deliberately preserving a small number of initial "sink" tokens in the attention window allows models trained with finite attention windows to generalize stably to much longer, effectively unbounded input streams without fine-tuning, by maintaining a more normal attention-score distribution. Related work on sparse and structured attention seeks to reduce the quadratic cost of full attention while preserving access to distant but relevant tokens, an important complement to purely retrieval-based mitigation because it addresses the underlying computational constraint rather than only the symptom of positional bias.

4.3 Positional-Encoding Extensions

A third family targets the extrapolation failure discussed in Section 3.2 directly, through techniques that adapt or extend positional encoding schemes to longer sequences than those seen in training. These include position interpolation and related rescaling methods for RoPE, and alternative encoding schemes such as ALiBi (Press et al., 2022) that were explicitly designed to extrapolate more gracefully than absolute encodings. Surveys of long-context transformer architecture (Huang et al., 2023; Liu, J. et al., 2025) catalogue an extensive and still-growing set of such extensions, reflecting substantial ongoing engineering investment in this dimension of the problem, even as evaluation results such as those from RULER (Hsieh et al., 2024) indicate that expanded nominal context length alone does not close the gap between claimed and effective context capacity.

4.4 Prompt Compression and Context Curation

A final class of mitigations operates upstream of the model, compressing or curating the input before it is presented for inference. These methods range from extractive or abstractive summarization of source documents prior to synthesis, to learned prompt-compression models that remove low-information tokens while preserving task-relevant content. Because such techniques reduce both the absolute length of the input and, potentially, the proportion of distractor material, they can mitigate both the positional and dilution mechanisms described in Section 3, though they introduce their own risk: compression that discards or paraphrases source content before synthesis can itself introduce factual drift, which is particularly consequential for research synthesis tasks where faithful attribution to specific sources is often a requirement rather than an incidental feature.

5. Research Challenges and Limitations

Despite substantial empirical documentation of context-window degradation, several limitations constrain what the existing literature can tell us about long-document research synthesis specifically.

Benchmark-task mismatch. Most degradation evidence derives from question-answering and synthetic retrieval tasks (needle-in-a-haystack, key-value retrieval) rather than from open-ended synthesis tasks that require comparing, reconciling, and aggregating claims across many sources. RULER's aggregation and tracing tasks are a partial exception, but even these remain synthetic and short of the open-ended, judgment-laden character of genuine research synthesis (Hsieh et al., 2024).

Inconsistent RAG-versus-long-context conclusions. As noted in Section 2.4, studies directly comparing retrieval-augmented and long-context approaches reach materially different conclusions depending on dataset construction, retriever quality, and whether questions answerable from parametric knowledge alone are filtered out (Li, Z. et al., 2024; Li, K. et al., 2025). This inconsistency limits practical guidance for designing synthesis pipelines and suggests that architecture choice cannot yet be generalized across task types without task-specific evaluation.

Absence of faithfulness-specific metrics. Accuracy-based benchmarks capture whether a correct fact was retrieved but not whether a synthesized summary faithfully represents the balance of evidence across sources, correctly attributes claims, or avoids conflating similar but distinct findings—failure modes especially relevant to context degradation, since a model that underweights middle-positioned sources may produce a synthesis that is fluent and superficially accurate on spot-checked facts while systematically misrepresenting the overall body of evidence.

Confounding of length and complexity. Because longer contexts in realistic settings also tend to contain more distractor content, more topics, and more potential for cross-document contradiction, it is difficult to cleanly separate degradation attributable to sequence length per se from degradation attributable to increased task complexity, a distinction that matters for choosing between architectural fixes (which primarily address length) and retrieval or curation fixes (which primarily address complexity and relevance density).

Rapid model turnover. Because context-length capabilities and mitigation techniques are advancing quickly, findings regarding a specific model's degradation profile risk becoming outdated; a recent case study found that at least one contemporary model showed little evidence of the lost-in-the-middle effect on simple factoid retrieval even near its context limit, suggesting that some mitigation strategies developed against earlier models may already be partially obsolete for certain tasks, even as the underlying aggregation and dilution mechanisms documented by more comprehensive benchmarks such as RULER appear to persist across model generations (Hsieh et al., 2024).

6. Research Gaps and Future Directions

Building on the limitations identified above, several directions merit priority attention from the research community.

6.1 Synthesis-Specific Evaluation Protocols

There is a clear need for benchmarks that evaluate long-document synthesis directly, rather than inferring synthesis capability from single-fact retrieval or question-answering accuracy. Such benchmarks would need to construct source-document sets in which claims of known ground-truth relevance are deliberately distributed across positions and documents, allowing researchers to measure not only whether a synthesized output is factually correct but whether its treatment of evidence is position-invariant—that is, whether the model's synthesis would remain materially the same if the order of input documents were permuted. This kind of permutation-invariance testing, a natural extension of the position-controlled methodology introduced by Liu et al. (2024), has not yet been systematically applied to open-ended synthesis outputs.

6.2 Faithfulness and Attribution Metrics for Long-Form Synthesis

Complementary to evaluation protocol design is the development of automated or semi-automated metrics capable of detecting under-attribution or over-attribution to particular sources in a synthesized document, distinguishing genuine claim aggregation from superficial paraphrase, and flagging cases where a synthesis omits a minority or dissenting finding that was nonetheless present in the source material. Existing faithfulness metrics developed for single-document summarization are unlikely to transfer cleanly to the multi-source case, since the relevant failure mode in synthesis is often selective omission across sources rather than local hallucination within a single passage.

6.3 Adaptive, Task-Aware Hybrid Architectures

Given the inconsistent comparative findings between RAG and long-context processing (Section 2.4), a promising direction is the further development of adaptive architectures that route or restructure processing based on the specific demands of a synthesis query—for example, detecting when a query requires cross-document comparison, which favors long-context or long-retrieval-unit processing, versus when it requires locating a small number of specific facts, which favors conventional retrieval. Extending Self-Route-style confidence-based routing (Li, Z. et al., 2024) to multi-step synthesis workflows, where different stages of a single task might benefit from different architectures, remains largely unexplored.

6.4 Mechanistic Understanding of Aggregation Failure

While positional bias and attention-sink phenomena are increasingly well characterized mechanistically, the aggregation and dilution failures most relevant to synthesis—why models undercount or fail to integrate facts distributed across many positions—remain comparatively underexplored at a mechanistic level. Interpretability work that traces how attention and representation patterns evolve as the number of relevant facts to be aggregated increases could inform both better training objectives and better architectural interventions targeted specifically at synthesis-style aggregation rather than single-fact retrieval.

6.5 Human-Comparable Synthesis Baselines

Finally, the field would benefit from studies that directly compare LLM-based long-document synthesis against expert human synthesis (for example, systematic-review methodology) on the same source sets, using the same evaluation criteria. Such comparisons would clarify not merely whether context degradation exists in a statistical sense, but whether its practical magnitude is comparable to, smaller than, or larger than the biases and errors already known to affect human research synthesis, providing a more actionable benchmark for deciding when LLM-assisted synthesis is reliable enough for a given use case.

7. Conclusion

Context-window degradation is a well-documented and multi-causal phenomenon: primacy-recency bias, positional-encoding extrapolation failure, attention dilution, and distractor accumulation each contribute to a growing gap between the nominal and effective context capacity of large language models. For long-document research synthesis, a task that structurally requires uniform, position-independent attention to many dispersed source passages, this gap is especially consequential, since it can produce synthesized outputs that are fluent and locally accurate while systematically misrepresenting the balance of evidence across sources. Current mitigation strategies—retrieval augmentation, hybrid routing, attention-sink management, positional-encoding extensions, and prompt compression—each address some but not all of the underlying mechanisms, and comparative evidence on their relative effectiveness remains inconsistent across studies and task types. Closing this gap will likely require synthesis-specific evaluation protocols, faithfulness and attribution metrics tailored to multi-source aggregation, and adaptive architectures capable of matching processing strategy to the specific demands of a given synthesis task, rather than relying on context-window size alone as a proxy for synthesis reliability.

Complete References

Bai, Y., Lv, X., Zhang, J., Lyu, H., Tang, J., Huang, Z., Du, Z., Liu, X., Zeng, A., Hou, L., Dong, Y., Tang, J., & Li, J. (2024). LongBench: A bilingual, multitask benchmark for long context understanding. In Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pp. 3119–3137. Association for Computational Linguistics. arXiv:2308.14508. https://doi.org/10.48550/arXiv.2308.14508

Hsieh, C.-P., Sun, S., Kriman, S., Acharya, S., Rekesh, D., Jia, F., Zhang, Y., & Ginsburg, B. (2024). RULER: What's the real context size of your long-context language models? In Proceedings of the First Conference on Language Modeling (COLM 2024). arXiv:2404.06654. https://doi.org/10.48550/arXiv.2404.06654

Huang, Y., Xu, J., Jiang, Z., Lai, J., Li, Z., Yao, Y., Chen, T., Yang, L., Xin, Z., & Ma, X. (2023). Advancing transformer architecture in long-context large language models: A comprehensive survey. arXiv:2311.12351. https://doi.org/10.48550/arXiv.2311.12351

Kamradt, G. (2023). Needle in a haystack: Pressure testing LLMs [Software repository]. GitHub. https://github.com/gkamradt/LLMTest_NeedleInAHaystack

Kazemnejad, A., Padhi, I., Ramamurthy, K. N., Das, P., & Reddy, S. (2023). The impact of positional encoding on length generalization in transformers. In Advances in Neural Information Processing Systems 36 (NeurIPS 2023). arXiv:2305.19466. https://doi.org/10.48550/arXiv.2305.19466

Li, K., Zhang, L., Jiang, Y., Xie, P., Huang, F., Wang, S., & Cheng, M. (2025). LaRA: Benchmarking retrieval-augmented generation and long-context LLMs — No silver bullet for LC or RAG routing. arXiv:2502.09977. https://doi.org/10.48550/arXiv.2502.09977

Li, Z., Li, C., Zhang, M., Mei, Q., & Bendersky, M. (2024). Retrieval augmented generation or long-context LLMs? A comprehensive study and hybrid approach. In Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: Industry Track (EMNLP 2024), pp. 881–893. Association for Computational Linguistics. arXiv:2407.16833. https://doi.org/10.48550/arXiv.2407.16833

Liu, J., Zhu, D., Bai, Z., He, Y., Liao, H., Que, H., Wang, Z., Zhang, C., Zhang, G., Zhang, J., Zhang, Y., Chen, Z., Guo, H., Li, S., Liu, Z., Shan, Y., Song, Y., Tian, J., Wu, W., Zhou, Z., Zhu, R., Feng, J., Gao, Y., He, S., Li, Z., Liu, T., Meng, F., Su, W., Tan, Y., Wang, Z., Yang, J., Ye, W., Zheng, B., Zhou, W., Huang, W., Li, S., & Zhang, Z. (2025). A comprehensive survey on long context language modeling. arXiv:2503.17407. https://doi.org/10.48550/arXiv.2503.17407

Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., & Liang, P. (2024). Lost in the middle: How language models use long contexts. Transactions of the Association for Computational Linguistics, 12, 157–173. https://doi.org/10.1162/tacl_a_00638 (arXiv:2307.03172)

Press, O., Smith, N. A., & Lewis, M. (2022). Train short, test long: Attention with linear biases enables input length extrapolation. In Proceedings of the Tenth International Conference on Learning Representations (ICLR 2022). arXiv:2108.12409. https://doi.org/10.48550/arXiv.2108.12409

Su, J., Ahmed, M., Lu, Y., Pan, S., Bo, W., & Liu, Y. (2024). RoFormer: Enhanced transformer with rotary position embedding. Neurocomputing, 568, 127063. https://doi.org/10.1016/j.neucom.2023.127063 (arXiv:2104.09864)

Xiao, G., Tian, Y., Chen, B., Han, S., & Lewis, M. (2024). Efficient streaming language models with attention sinks. In Proceedings of the Twelfth International Conference on Learning Representations (ICLR 2024). arXiv:2309.17453. https://doi.org/10.48550/arXiv.2309.17453

Xu, P., Ping, W., Wu, X., McAfee, L., Zhu, C., Liu, Z., Subramanian, S., Bakhturina, E., Shoeybi, M., & Catanzaro, B. (2024). Retrieval meets long context large language models. In Proceedings of the Twelfth International Conference on Learning Representations (ICLR 2024). arXiv:2310.03025. https://doi.org/10.48550/arXiv.2310.03025